



Lecture Notes on Pattern Recognition

Session 06(A): Discriminant Functions (I)

Road map

Discriminant Functions

- Discriminant
- Learning Discriminants
- Iterative Optimization
- Criterion Functions
- MSE Procedure

Statistical vs Discriminant

Linear Discriminants

Solution Space

Gradient Descent

Newton's Method

Perceptron Criterion

Relaxation Procedure

LMS

TULIP Team for Universal Learning and Intelligent Processing

FLIP: Pattern Recognition (G. Li)

DEAKIN
Worldly





Statistical vs Discriminant Approach

- **Statistical** approaches find the decision boundary by first estimating the probability distribution of the patterns belonging to each class.
- In the **discriminant** approach, the decision boundary is constructed explicitly without assuming a probability distribution.



DEAKIN
Worldly

Discriminant Approach

- Specify **parametric form** of the decision boundary (e.g., linear or quadratic).
- Find the **best** decision boundary of the specified form using a set of training examples x_k .
 - This is performed by **minimizing** a criterion function (e.g., “training error” or “sample risk”):

$$J(w) = \frac{1}{n} \sum_{k=1}^n [z_k - g(x_k, w)]^2$$

correct class predicted class

 **TULIP** (Tool for Universal Learning and Intelligent Processing)

FLJP: Pattern Recognition (Ch 1)

5

Discriminant Functions: two-categories case

Linear Discriminant Functions: two-categories case

- A linear discriminant function has the following form:
$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0 = \sum_{i=1}^d w_i x_i + w_0$$
- The decision boundary, given by $g(\mathbf{x})=0$, is a **hyperplane** where the orientation of the hyperplane is determined by \mathbf{w} and its location by w_0 .
 - \mathbf{w} is the normal to the hyperplane
 - If $w_0=0$, the hyperplane passes through the origin

 TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 7

Geometric Interpretation of $g(\mathbf{x})$

- $g(\mathbf{x})$ provides **an algebraic measure of the distance** of \mathbf{x} from the hyperplane.

\mathbf{x} can be expressed as follows:

$$\mathbf{x} = \mathbf{x}_p + r \frac{\mathbf{w}}{\|\mathbf{w}\|}$$

$$g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0 = \mathbf{w}^T (\mathbf{x}_p + r \frac{\mathbf{w}}{\|\mathbf{w}\|}) + w_0$$

$$= \mathbf{w}^T \mathbf{x}_p + r \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|} + w_0 = r \|\mathbf{w}\|$$

 TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 8

Geometric Interpretation of $g(\mathbf{x})$

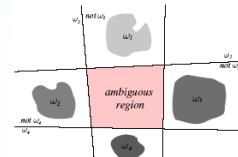
- Therefore, the distance of \mathbf{x} from the hyperplane is given by:
$$r = \frac{g(\mathbf{x})}{\|\mathbf{w}\|}$$
- w_0 determines the distance of the hyperplane from the origin:
 - setting $\mathbf{x}=0$: $r = \frac{w_0}{\|\mathbf{w}\|}$

 TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 9

Linear Discriminant Functions: multi-category case

- There are several ways to devise multi-category classifiers using linear discriminant functions:
 - One against the rest** (i.e., c-1 two-class problems)

problem • ambiguous regions

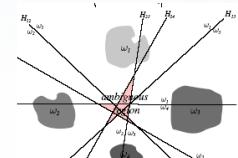


 TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 10

Linear Discriminant Functions: multi-category case

- There are several ways to devise multi-category classifiers using linear discriminant functions:
 - One against another** (i.e., c(c-1)/2 pairs of classes)

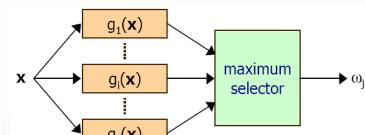
problem • ambiguous regions



 TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 11

Linear Discriminant Functions: multi-category case

- To avoid the problem of ambiguous regions:
 - Define c linear discriminant functions
 - Assign \mathbf{x} to ω_i if $g_i(\mathbf{x}) > g_j(\mathbf{x})$ for all $j \neq i$.
- The resulting classifier is called a **linear machine**



 TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 12

Linear Discriminant Functions: multi-category case

- A linear machine divides the feature space in c **convex** decisions regions.
 - If x is in region R_i , the $g_i(x)$ is the largest.

\bullet $c(c-1)/2$ pairs of regions but typically less decision boundaries

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 18

Linear Discriminant Functions: multi-category case

- The boundary between two regions R_i and R_j is a portion of the hyperplane H_{ij} given by:

$$g_i(\mathbf{x}) = g_j(\mathbf{x}) \text{ or } g_i(\mathbf{x}) - g_j(\mathbf{x}) = 0$$

$$\text{or } (\mathbf{w}_i - \mathbf{w}_j)' \mathbf{x} + (w_{i0} - w_{j0}) = 0$$
- $(\mathbf{w}_i - \mathbf{w}_j)$ is normal to H_{ij} and the signed distance from x to H_{ij} is

$$r = \frac{g_i(\mathbf{x}) - g_j(\mathbf{x})}{\|\mathbf{w}_i - \mathbf{w}_j\|}$$

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 19

Higher Order Discriminant Functions

- Can produce more complicated decision boundaries than linear discriminant functions.

Linear discriminant: $g(\mathbf{x}) = \mathbf{w}' \mathbf{x} + w_0 = \sum_{i=1}^d w_i x_i + w_0$

Quadratic discriminant: obtained by adding terms corresponding to products of pairs of components of x

$$g(\mathbf{x}) = w_0 + \sum_{i=1}^d w_i x_i + \sum_{i=1}^d \sum_{j=1}^d x_i x_j w_{ij}$$

Polynomial discriminant: obtained by adding terms such as $x_i x_j x_k w_{ijk}$.

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 15

Generalized discriminants

- Generalized linear discriminant functions

$$g(\mathbf{x}) = \sum_{i=1}^{\hat{d}} a_i y_i(\mathbf{x}) \text{ or } g(\mathbf{x}) = \mathbf{a}' \mathbf{y}$$
 - a is a \hat{d} -dimensional weight vector
 - y functions $y_i(\mathbf{x})$ map points from the d -dimensional \mathbf{x} -space to the \hat{d} -dimensional \mathbf{y} -space
 - usually $\hat{d} >> d$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} \Rightarrow \begin{bmatrix} y_1(\mathbf{x}) \\ y_2(\mathbf{x}) \\ \vdots \\ y_{\hat{d}}(\mathbf{x}) \end{bmatrix}$$

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 16

Generalized discriminants

- Generalized linear discriminant functions

$$g(\mathbf{x}) = \sum_{i=1}^{\hat{d}} a_i y_i(\mathbf{x}) \text{ or } g(\mathbf{x}) = \mathbf{a}' \mathbf{y}$$
 - The resulting discriminant function is **not** linear in x but it **is** linear in y .
 - The generalized discriminant separates points in the transformed space by a **hyperplane** passing through the **origin**.

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 17

Generalized discriminants

- Example

$$g(\mathbf{x}) = -1 + x + 2x^2$$

$$\mathbf{a} = \begin{bmatrix} -1 \\ 1 \\ 2 \end{bmatrix}$$

$$\mathbf{y} = \begin{bmatrix} y_1(\mathbf{x}) \\ y_2(\mathbf{x}) \\ y_3(\mathbf{x}) \end{bmatrix} = \begin{bmatrix} 1 \\ x \\ x^2 \end{bmatrix}$$
 - The corresponding decision regions R_1, R_2 in the \mathbf{x} -space are **not** simply connected!
 - $g(\mathbf{x})$ maps a line in \mathbf{x} -space to a parabola in \mathbf{y} -space.
 - The plane $\mathbf{a}'\mathbf{y}=0$ divides the \mathbf{y} -space in two decision regions \hat{R}_1, \hat{R}_2 .

$g(x) > 0 \text{ if } x < -1 \text{ or } x > 0.5$

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 18

Learning Discriminants



- Linear Separable
- Solution Space

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 19

Learning: two-category, linearly separable case

- Task: Given a linear discriminant function, **learn** the weights \mathbf{w} using a set of n labeled samples \mathbf{x}_i
 - i.e., each \mathbf{x}_i has a class label ω_1 or ω_2

$$g(\mathbf{x}) = \mathbf{w}' \mathbf{x} + w_0 = \sum_{i=1}^d w_i x_i + w_0$$

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 20

Augmented feature/weight vectors

- Increasing dimensionality simplifies things:
 - Find a single weight vector \mathbf{a} in \mathbf{y} -space.
 - Decision hyperplane, given by $g(\mathbf{y})=0$, in \mathbf{y} -space passes through the origin (but could be at any position in \mathbf{x} -space).

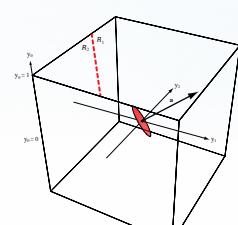
$$g(\mathbf{x}) = \mathbf{w}' \mathbf{x} + w_0 = \sum_{i=1}^d w_i x_i + x_0 w_0 = \sum_{i=0}^d w_i x_i = \mathbf{a}' \mathbf{y}$$

$$\mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \\ \dots \\ w_d \end{bmatrix} \Rightarrow \mathbf{a} = \begin{bmatrix} w_0 \\ w_1 \\ \dots \\ w_d \end{bmatrix} \quad \mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \dots \\ x_d \end{bmatrix} \Rightarrow \mathbf{y} = \begin{bmatrix} 1 \\ x_1 \\ x_2 \\ \dots \\ x_d \end{bmatrix} \quad (\text{i.e., } x_0 = 1)$$

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 21

Classification in augmented space

- Given a discriminant function $g(\mathbf{x})=\mathbf{a}'\mathbf{y}$, the classification rule is:
 - If $\mathbf{a}'\mathbf{y}_i > 0$ assign \mathbf{y}_i to ω_1
 - else if $\mathbf{a}'\mathbf{y}_i < 0$ assign \mathbf{y}_i to ω_2



TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 22

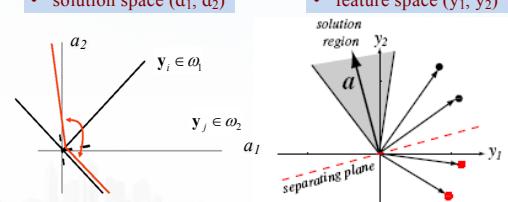
Learning: two-category, linearly separable case

- Given a discriminant function $g(\mathbf{x})=\mathbf{a}'\mathbf{y}$, the goal is to **learn** the weights \mathbf{a} using a set of n labeled samples \mathbf{y}_i
 - Every training sample \mathbf{y}_i places a constraint on the weight vector \mathbf{a} .
 - $\mathbf{a}'\mathbf{y}=0$ defines a hyperplane in solution space having \mathbf{y} as a normal vector.
 - Given n examples, the solution \mathbf{a} must lie on the intersection of n half-spaces

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 23

Effect of training examples on solution

- Solution can be visualized either in the solution space or the feature space
 - solution space (a_1, a_2)
 - feature space (y_1, y_2)



TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 24

Solution Uniqueness/Constraints

- Solution vector is usually not unique; we can impose certain constraints, e.g.:
 - Method 1:** find unit-length weight vector that maximizes the minimum distance from the samples to the separating plane.

FLIP: Pattern Recognition (G. Li) 25

Solution Uniqueness/Constraints

- Solution vector is usually not unique; we can impose certain constraints, e.g.:
 - Method 2:** find minimum-length weight vector satisfying the constraint below where b is a positive constant called **margin**.

$$g(\mathbf{x}) = \mathbf{a}^T \mathbf{y}_i \geq b > 0$$

Main objective:

- move solution to the center of the feasible region as this solution is more likely to classify new test samples correctly.

FLIP: Pattern Recognition (G. Li) 26

Normalized Version

- If \mathbf{y}_i in ω_2 , replace \mathbf{y}_i by $-\mathbf{y}_i$
- Find \mathbf{a} such that: $\mathbf{a}^T \mathbf{y}_i > 0$

FLIP: Pattern Recognition (G. Li) 27

Iterative Optimization

- Gradient Descent
- Newton's Method

FLIP: Pattern Recognition (G. Li) 28

Iterative Optimization

- Define an error function $J(\mathbf{a})$ (based on training samples) that is minimized if \mathbf{a} is a solution vector.
- Minimize $J(\mathbf{a})$ iteratively:

$$\mathbf{a}(k+1) = \mathbf{a}(k) + \eta(k) \mathbf{p}_k$$

FLIP: Pattern Recognition (G. Li) 29

Gradient Descent

- Gradient descent rule:

$$\mathbf{p}_k = -\nabla J(\mathbf{a}(k))$$

Algorithm 1 (Basic gradient descent)

```

1 begin initialize a, criterion θ, η(·), k = 0
2 do k ← k + 1
3   a ← a - η(k) ∇ J(a)
4   until η(k) ∇ J(a) < θ
5 return a
6 end

```

$$\mathbf{a}(k+1) = \mathbf{a}(k) + \eta(k) \mathbf{p}_k$$

FLIP: Pattern Recognition (G. Li) 30

Gradient Descent

- solution space \mathbf{a}

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 31

Gradient Descent

- What is the effect of the learning rate?

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 32

- slow but converges to solution
- fast by overshoots solution

Gradient Descent

- How to choose the learning rate $\eta(k)$?
 - Make sure $J(\mathbf{a}(k+1))$ reaches the lowest point along the decreasing direction

Taylor series approximation

$$J(\mathbf{a}) \simeq J(\mathbf{a}(k)) + \nabla J^t(\mathbf{a} - \mathbf{a}(k)) + \frac{1}{2}(\mathbf{a} - \mathbf{a}(k))^t \mathbf{H} (\mathbf{a} - \mathbf{a}(k)),$$

Hessian (2nd derivatives)

optimum learning rate

If $J(\mathbf{a})$ is quadratic, then \mathbf{H} is constant which implies that the learning rate is constant!

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 33

Newton's method

- Newton's descent rule:

```

1 begin initialize a, criterion θ
2   do
3     a ← a - H-1∇J(a)
4     until H-1∇J(a) < θ
5   return a
6 end

```

$\mathbf{p}_k = -\mathbf{H}^{-1}\nabla J(\mathbf{a}(k))$

requires inverting \mathbf{H}

$\mathbf{a}(k+1) = \mathbf{a}(k) + \eta(k)\mathbf{p}_k$

\mathbf{p}_k • search direction

$\mathbf{a}(k) \xrightarrow{\eta(k)\mathbf{p}_k} \mathbf{a}(k+1)$

$\eta(k)$ • learning rate

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 34

Newton's method

- If $J(\mathbf{a})$ is quadratic, Newton's method converges in **one step**!

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 35

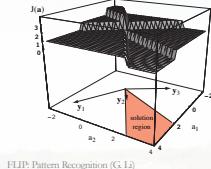
Criterion Function

- Perceptron Criterion Function
- Relaxation Procedure

TULIP Team for Universal Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 36

Criterion Function

- Constructing a criterion function for solving the linear inequalities $\mathbf{a}'\mathbf{y}_i > 0$.
 - The simplest choice is let $J(\mathbf{a})$ be the number of samples misclassified by \mathbf{a} .
 - This leads to a piecewise constant, not suitable for gradient search



FLIP: Pattern Recognition (G. 1j)

37



Perceptron Criterion Function

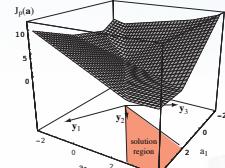
- Constructing a criterion function for solving the linear inequalities $\mathbf{a}'\mathbf{y}_i > 0$.

Perceptron Criterion Function:

$$J_p(\mathbf{a}) = \sum_{\mathbf{y} \in Y(\mathbf{a})} (-\mathbf{a}' \mathbf{y})$$

where $Y(\mathbf{a})$ is the set of samples misclassified by \mathbf{a} .

If $Y(\mathbf{a})$ is empty, $J_p(\mathbf{a})=0$; otherwise, $J_p(\mathbf{a}) > 0$ (normalized version)



FLIP: Pattern Recognition (G. 1j)

38

Perceptron rule (Batch Correction)

- The gradient of $J_p(\mathbf{a})$ is:

$$\nabla J_p = \sum_{\mathbf{y} \in Y(\mathbf{a})} (-\mathbf{y})$$

- The perceptron update rule is obtained using gradient descent:

$$\mathbf{a}(k+1) = \mathbf{a}(k) - \eta(k) \nabla J(\mathbf{a}(k)).$$

$$\mathbf{a}(k+1) = \mathbf{a}(k) + \eta(k) \sum_{\mathbf{y} \in Y(\mathbf{a})} \mathbf{y}$$



FLIP: Pattern Recognition (G. 1j)

39



Perceptron rule (Single-Sample Correction)

- Fix the learning rate $\eta(k) = 1$
- Modify the weight vector upon each misclassified single sample
- The misclassified sample will be input infinitely often
- The perceptron update rule is obtained using gradient descent:

$$\mathbf{a}(k+1) = \mathbf{a}(k) + \eta(k) \mathbf{y}$$



FLIP: Pattern Recognition (G. 1j)

40

Perceptron rule (Single-Sample Correction)

Perceptron Convergence Theorem:

- If training samples are linearly separable, then the sequence of weight vectors will terminate at a solution vector in a finite number of steps.

Algorithm 4 (Fixed-increment single-sample Perceptron)

```

1 begin initialize a, k = 0
2   do k ← (k + 1)modn    η(k)=1
3     if y_k is misclassified by a then a ← a + y_k
4     until all patterns properly classified
5   return a
6 end
  • consider one example at a time

```

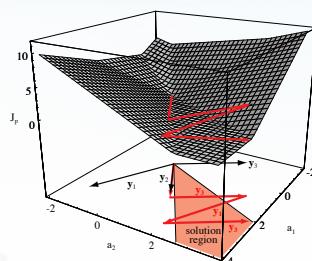


FLIP: Pattern Recognition (G. 1j)

41



Perceptron rule



FLIP: Pattern Recognition (G. 1j)

42

Perceptron rule (Margins)

- Generalizations:
 - Variable learning rate and a margin

Algorithm 5 (Variable increment Perceptron with margin)

```

1 begin initialize a, criterion  $\theta$ , margin b,  $\eta(\cdot)$ , k = 0
2   do k ← k + 1
3     if  $a^t y_k + b < 0$  then a ← a +  $\eta(k) y_k$ 
4     until  $a^t y_k + b \leq 0$  for all k
5   return a
6 end
```



FLIP: Pattern Recognition (G. Li)

43

Relaxation Procedures

- Note that different criterion functions exist
- One possible choice is:

$$J_q(a) = \sum_{y \in Y} (a^t y)^2$$

Where Y is again the set of the training samples that are misclassified by a

- However, there are two problems with this criterion
 - The function is too smooth and can converge to a=0
 - J_q is dominated by training samples with large magnitude



FLIP: Pattern Recognition (G. Li)

44

Relaxation Procedures

- A modified version that avoids two problems is

$$J_r(a) = \frac{1}{2} \sum_{y \in Y} \frac{(a^t y - b)^2}{\|y\|^2}$$

- Here Y is the set of samples for which its gradient is given by:

$$\nabla J_r = \sum_{y \in Y} \frac{(a^t y - b)y}{\|y\|^2}$$



FLIP: Pattern Recognition (G. Li)

45

Relaxation Procedures

Algorithm 8 (Batch relaxation with margin)

```

1 begin initialize a,  $\eta(\cdot)$ , k = 0
2   do k ← k + 1
3      $y_k = \{\}$ 
4     j = 0
5     do j ← j + 1
6       if  $y_j$  is misclassified then Append  $y_j$  to  $y_k$ 
7     until j = n
8     a ← a +  $\eta(k) \sum_{y \in y_k} \frac{b - a^t y}{\|y\|^2} y$ 
9   until  $y_k = \{\}$ 
10 return a
11 end
```



FLIP: Pattern Recognition (G. Li)

46

Relaxation Procedures

Algorithm 9 (Single-sample relaxation with margin)

```

1 begin initialize a,  $\eta(\cdot)$ , k = 0
2   do k ← k + 1
3     if  $y_k$  is misclassified then a ← a +  $\eta(k) \frac{b - a^t y}{\|y_k\|^2} y_k$ 
4     until all patterns properly classified
5   return a
6 end
```

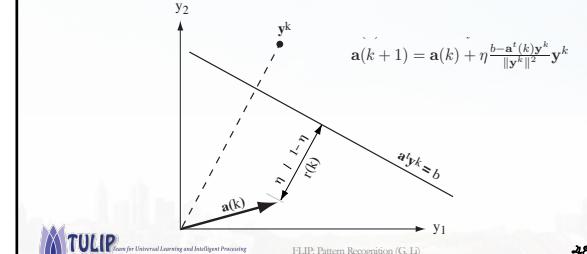


FLIP: Pattern Recognition (G. Li)

47

Relaxation Procedures

- Since $y / \|y\|$ is the unit normal vector for the hyperplane. The correction calls for a(k) to be moved a certain fraction of η of the distance from a(k) to the hyperplane



FLIP: Pattern Recognition (G. Li)

48

Nonseparable Behavior

- Perceptron and relaxation procedures are all error-correction procedure.
 - Their success on separable problems is largely due to its relentless search for an error-free solution.
 - For nonseparable problems, correction may never converge.



FLIP: Pattern Recognition (G. Li)

49

MSE Procedure

- MSE and Discrimants
- Pseudoinverse



FLIP: Pattern Recognition (G. Li)

50

Minimum Squared Error Procedures

- Define a criterion function for all samples, rather than the misclassified samples
- Try to find the vector such that $\mathbf{a}^t \mathbf{y}_i > b_i$
 - Where b_i is any positive value

$$\begin{pmatrix} Y_{10} & Y_{11} & \cdots & Y_{1d} \\ Y_{20} & Y_{21} & \cdots & Y_{2d} \\ \vdots & \vdots & & \vdots \\ \vdots & \vdots & & \vdots \\ \vdots & \vdots & & \vdots \\ Y_{n0} & Y_{n1} & \cdots & Y_{nd} \end{pmatrix} \begin{pmatrix} a_0 \\ a_1 \\ \vdots \\ a_d \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{pmatrix}$$

FLIP: Pattern Recognition (G. Li)

51

Minimum Squared Error Procedures

- Minimum squared error and pseudoinverse
 - The problem is to find a weight vector \mathbf{a} satisfying $\mathbf{Y}\mathbf{a} = \mathbf{b}$
 - If we have more equations than unknowns, \mathbf{a} is over-determined.
 - We want to choose the one that minimizes the sum-of-squared-error criterion function

$$J_s(\mathbf{a}) = \|\mathbf{Y}\mathbf{a} - \mathbf{b}\|^2 = \sum_{i=1}^n (\mathbf{a}^t \mathbf{y}_i - b_i)^2.$$

FLIP: Pattern Recognition (G. Li)

52

Minimum Squared Error Procedures

- Pseudoinverse

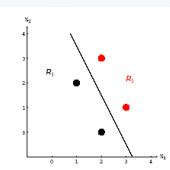
$$\begin{aligned} \mathbf{a} &= (\mathbf{Y}^t \mathbf{Y})^{-1} \mathbf{Y}^t \mathbf{b} \\ &= \mathbf{Y}^t \mathbf{b}, \end{aligned}$$

Algorithm 10 (LMS)

```

1 begin initialize a, b, criterion θ, η(·), k = 0
2   do k ← k + 1
3     a ← a + η(k)(b_k - a^t y^k) y^k
4   until η(k)(b_k - a^t y^k) y^k < θ
5   return a
6 end

```



FLIP: Pattern Recognition (G. Li)

53

Seminar S09

- Topics**
 - Feature Selection and Feature Extraction
 - From Perceptron to Neural Network
 - Widrow-Hoff Algorithm and LMS
- Requirements**
 - Prepare a **15 minutes** talk on your chosen topic
 - Make **ppt** to assist your talk
 - Prepare **at least 3 questions** to ask the audience after your talk
 - Get ready to **take questions** from the audience
- Hints**
 - You can search for research articles from Google Scholar



FLIP: Pattern Recognition (G. Li)

54

