



Lecture Notes on
Pattern Recognition

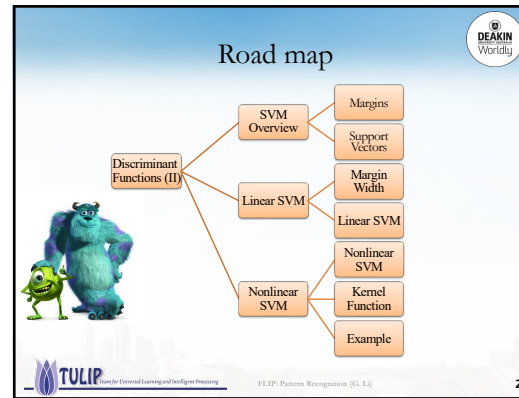
Session 07(B): Discriminant Functions (II)

Gang Li
School of Information Technology
Deakin University, VIC 3125, Australia

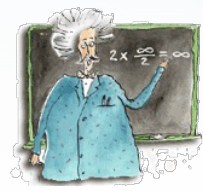
FLIP: Pattern Recognition (G. Li)

1





2

SVM Overview



- Best Linear Classifiers
- Margins
- Support Vectors
- VC Dimension

FLIP: Pattern Recognition (G. Li)



3

Learning through
empirical risk minimization

- Estimate $g(x)$ from a finite set of observations by minimizing some kind of error function, for example, the **empirical risk (sample risk)**:

$$R_{emp}(w, w_0) = \frac{1}{n} \sum_{k=1}^n [z_k - g(x_k, w, w_0)]^2$$

class labels: $z_k = \begin{cases} +1 & \text{if } \mathbf{x}_k \in \omega_1 \\ -1 & \text{if } \mathbf{x}_k \in \omega_2 \end{cases}$






FLIP: Pattern Recognition (G. Li)

4

Learning through
empirical risk minimization

- Conventional **empirical risk** minimization over the training data **does not** imply good generalization to novel test data.
 - There could be a number of different functions which all approximate the training data set well.
 - Difficult to determine a function which **best** captures the true underlying structure of the data distribution
 - i.e., has good generalization capabilities

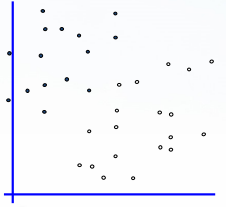



FLIP: Pattern Recognition (G. Li)

5



Linear Classifiers

• denotes +1
• denotes -1




$f(\mathbf{x}; \mathbf{w}, b) = \text{sign}(\mathbf{w} \cdot \mathbf{x} - b)$

How would you classify this data?





FLIP: Pattern Recognition (G. Li)


6




7




8




9




10




11



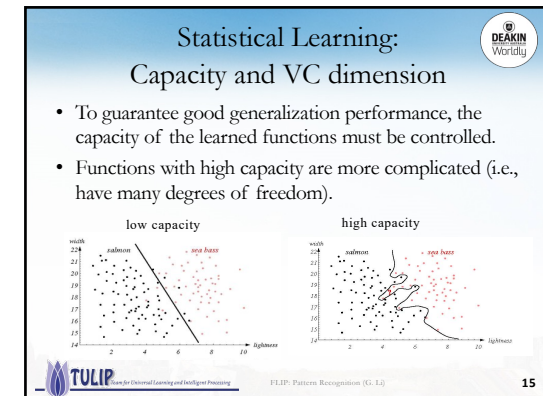
12



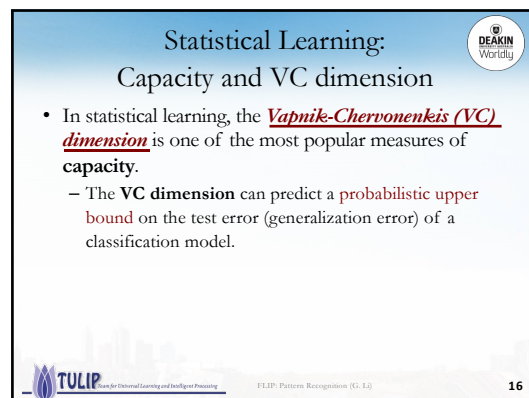
13



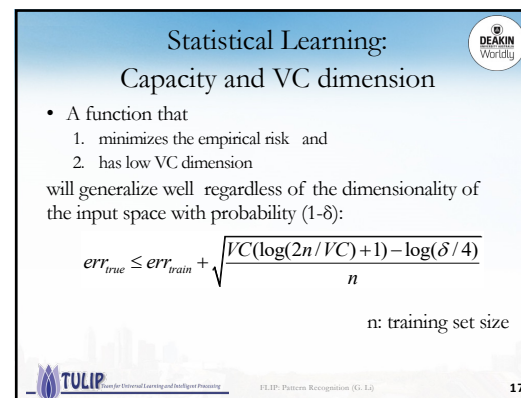
14



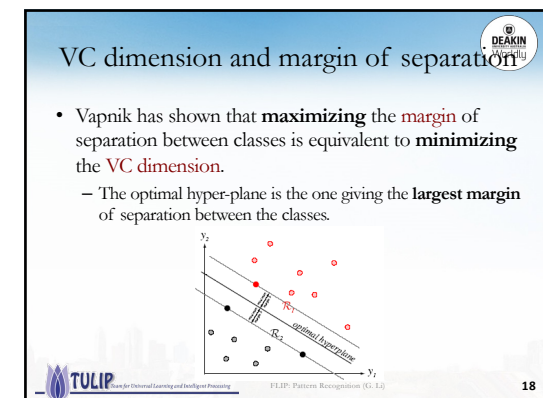
15



16



17



18

Margin of separation and support vectors

- The margin (i.e., empty area around the decision boundary) is defined by the distance to the nearest training patterns which we refer to as **support vectors**.
 - Intuitively speaking, these are the most difficult patterns to classify.

TULIP TU/e University of Twente Learning and Intelligent Processing FLIP: Pattern Recognition (G.1.1)

19

Margin of separation and support vectors

TULIP TU/e University of Twente Learning and Intelligent Processing FLIP: Pattern Recognition (G.1.1)

20

SVM Overview

- SVMs perform **structural risk minimization** to achieve good generalization performance.
- The optimization criterion is the **margin** of separation between classes.
- Training is equivalent to solving a **quadratic programming** problem with **linear constraints**.
- Primarily **two-class** classifiers but can be extended to **multiple classes**.

TULIP TU/e University of Twente Learning and Intelligent Processing FLIP: Pattern Recognition (G.1.1)

21

Linear SVM

- Margin Width
- Linear SVM

TULIP TU/e University of Twente Learning and Intelligent Processing FLIP: Pattern Recognition (G.1.1)

22

Specifying a line and margin

- How do we represent this mathematically?
 - ...in m input dimensions?

TULIP TU/e University of Twente Learning and Intelligent Processing FLIP: Pattern Recognition (G.1.1)

23

Specifying a line and margin

- Plus-plane = $\{ \mathbf{x} : \mathbf{w} \cdot \mathbf{x} + w_0 = +1 \}$
- Minus-plane = $\{ \mathbf{x} : \mathbf{w} \cdot \mathbf{x} + w_0 = -1 \}$

Classify as...

+	if $\mathbf{w} \cdot \mathbf{x} + w_0 \geq 1$
-	if $\mathbf{w} \cdot \mathbf{x} + w_0 \leq -1$
0	if $-1 < \mathbf{w} \cdot \mathbf{x} + w_0 < 1$

TULIP TU/e University of Twente Learning and Intelligent Processing FLIP: Pattern Recognition (G.1.1)

24

Specifying a line and margin

- Plus-plane = $\{ \mathbf{x} : \mathbf{w} \cdot \mathbf{x} + w_0 = +1 \}$
- Minus-plane = $\{ \mathbf{x} : \mathbf{w} \cdot \mathbf{x} + w_0 = -1 \}$
- Claim: The vector \mathbf{w} is perpendicular to the Plus Plane. Why?

TULIP: Towards Universal Learning and Intelligent Processing

FLIP: Pattern Recognition (G. Li)

25

Specifying a line and margin

- Plus-plane = $\{ \mathbf{x} : \mathbf{w} \cdot \mathbf{x} + w_0 = +1 \}$
- Minus-plane = $\{ \mathbf{x} : \mathbf{w} \cdot \mathbf{x} + w_0 = -1 \}$
- Claim: the vector \mathbf{w} is perpendicular to the Plus Plane. Why?

And so of course the vector \mathbf{w} is also perpendicular to the Minus Plane

Let \mathbf{u} and \mathbf{v} be two vectors on the Plus Plane. What is $\mathbf{w} \cdot (\mathbf{u} - \mathbf{v})$?

TULIP: Towards Universal Learning and Intelligent Processing

FLIP: Pattern Recognition (G. Li)

26

Specifying a line and margin

- Plus-plane = $\{ \mathbf{x} : \mathbf{w} \cdot \mathbf{x} + w_0 = +1 \}$
- Minus-plane = $\{ \mathbf{x} : \mathbf{w} \cdot \mathbf{x} + w_0 = -1 \}$
- The vector \mathbf{w} is perpendicular to the Plus Plane.
 - Let \mathbf{x}^* be any point on the minus plane
 - Let \mathbf{x}^+ be the closest plus-plane-point to \mathbf{x}^* .

Any location in \mathbb{R}^n , not necessarily a data point

TULIP: Towards Universal Learning and Intelligent Processing

FLIP: Pattern Recognition (G. Li)

27

Specifying a line and margin

- Plus-plane = $\{ \mathbf{x} : \mathbf{w} \cdot \mathbf{x} + w_0 = +1 \}$
- Minus-plane = $\{ \mathbf{x} : \mathbf{w} \cdot \mathbf{x} + w_0 = -1 \}$
- The vector \mathbf{w} is perpendicular to the Plus Plane
- Claim: $\mathbf{x}^+ = \mathbf{x}^* + \lambda \mathbf{w}$ for some value of λ . Why?

TULIP: Towards Universal Learning and Intelligent Processing

FLIP: Pattern Recognition (G. Li)

28

Specifying a line and margin

- Plus-plane = $\{ \mathbf{x} : \mathbf{w} \cdot \mathbf{x} + w_0 = +1 \}$
- Minus-plane = $\{ \mathbf{x} : \mathbf{w} \cdot \mathbf{x} + w_0 = -1 \}$
- The vector \mathbf{w} is perpendicular to the Plus Plane
- Claim: $\mathbf{x}^+ = \mathbf{x}^* + \lambda \mathbf{w}$ for some value of λ . Why?

The line from \mathbf{x}^* to \mathbf{x}^+ is perpendicular to the planes.

So to get from \mathbf{x}^* to \mathbf{x}^+ travel some distance in direction \mathbf{w} .

TULIP: Towards Universal Learning and Intelligent Processing

FLIP: Pattern Recognition (G. Li)

29

Specifying a line and margin

- What we know:
 - $\mathbf{w} \cdot \mathbf{x}^+ + w_0 = +1$
 - $\mathbf{w} \cdot \mathbf{x}^* + w_0 = -1$
 - $\mathbf{x}^+ = \mathbf{x}^* + \lambda \mathbf{w}$
 - $|\mathbf{x}^+ - \mathbf{x}^*| = M$

$M = \text{Margin Width}$

How do we compute M in terms of \mathbf{w} and b ?

TULIP: Towards Universal Learning and Intelligent Processing

FLIP: Pattern Recognition (G. Li)

30

Specifying a line and margin

- What we know:
 - $\mathbf{w} \cdot \mathbf{x}^+ + w_0 = +1$
 - $\mathbf{w} \cdot \mathbf{x}^- + w_0 = -1$
 - $\mathbf{x}^+ = \mathbf{x}^- + \lambda \mathbf{w}$
 - $|\mathbf{x}^+ - \mathbf{x}^-| = M$

$$\mathbf{w} \cdot (\mathbf{x}^- + \lambda \mathbf{w}) + w_0 = 1$$

$$\Rightarrow \mathbf{w} \cdot \mathbf{x}^- + w_0 + \lambda \mathbf{w} \cdot \mathbf{w} = 1$$

$$\Rightarrow -1 + \lambda \mathbf{w} \cdot \mathbf{w} = 1$$

$$\Rightarrow \lambda = \frac{2}{\mathbf{w} \cdot \mathbf{w}}$$

$M = \text{Margin Width}$

- How do we compute M in terms of \mathbf{w} and b ?

31

Specifying a line and margin

- What we know:
 - $\mathbf{w} \cdot \mathbf{x}^+ + w_0 = +1$
 - $\mathbf{w} \cdot \mathbf{x}^- + w_0 = -1$
 - $\mathbf{x}^+ = \mathbf{x}^- + \lambda \mathbf{w}$
 - $|\mathbf{x}^+ - \mathbf{x}^-| = M$
 - $\lambda = \frac{2}{\mathbf{w} \cdot \mathbf{w}}$

$$M = |\mathbf{x}^+ - \mathbf{x}^-| = |\lambda \mathbf{w}|$$

$$= \lambda |\mathbf{w}| = \lambda \sqrt{\mathbf{w} \cdot \mathbf{w}}$$

$$= \frac{2\sqrt{\mathbf{w} \cdot \mathbf{w}}}{\mathbf{w} \cdot \mathbf{w}} = \frac{2}{\sqrt{\mathbf{w} \cdot \mathbf{w}}}$$

$M = \text{Margin Width}$

- How do we compute M in terms of \mathbf{w} and b ?

32

Specifying a line and margin

- Given a guess of \mathbf{w} and w_0 we can
 - Compute whether all data points in the correct half-planes
 - Compute the width of the margin
- So now we need to search the space of \mathbf{w} 's to find the widest margin that matches all the data points. **How?**

$M = \text{Margin Width} = \frac{2}{\sqrt{\mathbf{w} \cdot \mathbf{w}}}$

33

Linear SVM: separable case

- Linear discriminant

$$g(\mathbf{x}) = \mathbf{w}'\mathbf{x} + w_0$$

Decide ω_1 if $g(\mathbf{x}) > 0$ and ω_2 if $g(\mathbf{x}) < 0$
- Class labels

$$z_k = \begin{cases} +1 & \text{if } \mathbf{x}_k \in \omega_1 \\ -1 & \text{if } \mathbf{x}_k \in \omega_2 \end{cases}$$
- Normalized version

$$z_k g(\mathbf{x}_k) > 0 \quad \text{or} \quad z_k (\mathbf{w}'\mathbf{x}_k + w_0) > 0, \quad \text{for } k = 1, 2, \dots, n$$

34

Linear SVM: separable case

- The distance of a point \mathbf{x}_k from the separating hyper-plane should satisfy the **constraint**:

$$\frac{z_k g(\mathbf{x}_k)}{\|\mathbf{w}\|} \geq b, \quad b > 0$$
- To ensure uniqueness, impose:

$$b \|\mathbf{w}\| = 1$$
- The above **constraint** becomes:

$$z_k g(\mathbf{x}_k) \geq 1 \quad \text{or} \quad z_k (\mathbf{w}'\mathbf{x}_k + w_0) > 1 \quad \text{where } b = \frac{1}{\|\mathbf{w}\|}$$

35

Linear SVM: separable case

- Optimization:
 - Maximize the margin $\frac{2}{\sqrt{\mathbf{w} \cdot \mathbf{w}}}$
 - With constraints \Rightarrow **Problem 1:** Minimize $\frac{1}{2} \|\mathbf{w}\|^2$ subject to $z_k (\mathbf{w}'\mathbf{x}_k + w_0) \geq 1, \quad k = 1, 2, \dots, n$

- Quadratic programming problem!

36

Linear SVM: separable case

- Use Langrange optimization, we seek to minimize:

$$L(\mathbf{w}, w_0, \lambda) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{k=1}^n \lambda_k [z_k (\mathbf{w}' \mathbf{x}_k + w_0) - 1], \quad \lambda_k \geq 0$$

- Easier to solve the “dual” problem

– Kuhn-Tucker construction:

$$\text{Problem 2: Maximize } \sum_{k=1}^n \lambda_k - \frac{1}{2} \sum_{k,j} \lambda_k \lambda_j z_k z_j \mathbf{x}_k' \mathbf{x}_j$$

$$\text{subject to } \sum_{k=1}^n z_k \lambda_k = 0, \quad \lambda_k \geq 0, \quad k = 1, 2, \dots, n$$



FLIP: Pattern Recognition (G. Li)

37

37

Linear SVM: separable case

- The solution is given by:

$$\mathbf{w} = \sum_{k=1}^n z_k \lambda_k \mathbf{x}_k, \quad w_0 = z_k - \mathbf{w}' \mathbf{x}_k$$

$$g(\mathbf{x}) = \mathbf{w}' \mathbf{x} + w_0$$

$$g(\mathbf{x}) = \sum_{k=1}^n z_k \lambda_k (\mathbf{x}_k' \mathbf{x}) + w_0 = \sum_{k=1}^n z_k \lambda_k (\mathbf{x} \cdot \mathbf{x}_k) + w_0$$

• dot product

- It can be shown that if \mathbf{x}_k is not a support vector, then the corresponding $\lambda_k = 0$.
 - Only support vectors contribute to the solution!



FLIP: Pattern Recognition (G. Li)

38

38

Linear SVM: non-separable case

- Allow miss-classifications (i.e., soft margin classifier) by introducing positive error (slack) variables ψ_k :

$$z_k (\mathbf{w}' \mathbf{x}_k + w_0) \geq 1 - \psi_k, \quad \psi_k \geq 0, \quad k = 1, 2, \dots, n$$

$$\text{Problem 3: Minimize } \frac{1}{2} \|\mathbf{w}\|^2 + c \sum_{k=1}^n \psi_k$$

$$\text{subject to } z_k (\mathbf{w}' \mathbf{x}_k + w_0) \geq 1 - \psi_k, \quad k = 1, 2, \dots, n$$

- The result is a hyperplane that minimizes the sum of errors ψ_k while maximizing the margin for the correctly classified data.

- constant c controls the trade-off between the margin and misclassification errors.
- Aims to prevent outliers from affecting the optimal hyperplane.



FLIP: Pattern Recognition (G. Li)

39

39

Linear SVM: non-separable case

- Easier to solve the “dual” problem

– Kuhn-Tucker construction

$$\text{Problem 4: Maximize } \sum_{k=1}^n \lambda_k - \frac{1}{2} \sum_{k,j} \lambda_k \lambda_j z_k z_j \mathbf{x}_k' \mathbf{x}_j$$

$$\text{subject to } \sum_{k=1}^n z_k \lambda_k = 0 \text{ and } 0 \leq \lambda_k \leq c, \quad k = 1, 2, \dots, n$$

where the use of error variables ψ_k constraint the range of the Lagrange coefficients from 0 to c .



FLIP: Pattern Recognition (G. Li)

40

40

Nonlinear SVM



- Nonlinear SVM
- Kernel Functions
- Example



FLIP: Pattern Recognition (G. Li)

41

41

Nonlinear SVM

- Extending these concepts to the non-linear case involves mapping the data to a high-dimensional space Φ :

$$\mathbf{x}_k \rightarrow \Phi(\mathbf{x}_k) = \begin{bmatrix} \phi_1(\mathbf{x}_k) \\ \phi_2(\mathbf{x}_k) \\ \vdots \\ \phi_h(\mathbf{x}_k) \end{bmatrix}$$

- Mapping the data to a sufficiently high dimensional space is likely to cast the data linearly separable in that space.



FLIP: Pattern Recognition (G. Li)

42

42

Nonlinear SVM

• linear SVM:

$$g(\mathbf{x}) = \sum_{k=1}^n z_k \lambda_k (\mathbf{x} \cdot \mathbf{x}_k) + w_0$$

• non-linear SVM:

$$g(\mathbf{x}) = \sum_{k=1}^n z_k \lambda_k (\Phi(\mathbf{x}) \cdot \Phi(\mathbf{x}_k)) + w_0$$

• Decide w_1 if $g(\mathbf{x}) > 0$ and w_2 if $g(\mathbf{x}) < 0$

FLIP: Pattern Recognition (G.1)
43

43

Nonlinear SVM

- The **disadvantage** of this approach is that the mapping might be very computationally intensive to compute!

$$\mathbf{x}_k \rightarrow \Phi(\mathbf{x}_k)$$

• non-linear SVM:

$$g(\mathbf{x}) = \sum_{k=1}^n z_k \lambda_k (\Phi(\mathbf{x}) \cdot \Phi(\mathbf{x}_k)) + w_0$$

FLIP: Pattern Recognition (G.1)
44

44

Nonlinear SVM

- The kernel trick**
 - Compute dot products using a kernel function

$$K(\mathbf{x}, \mathbf{x}_k) = \Phi(\mathbf{x}) \cdot \Phi(\mathbf{x}_k)$$

$$g(\mathbf{x}) = \sum_{k=1}^n z_k \lambda_k (\Phi(\mathbf{x}) \cdot \Phi(\mathbf{x}_k)) + w_0$$
$$g(\mathbf{x}) = \sum_{k=1}^n z_k \lambda_k K(\mathbf{x}, \mathbf{x}_k) + w_0$$

FLIP: Pattern Recognition (G.1)
45

45

Nonlinear SVM

- The kernel trick**
 - Compute dot products using a kernel function
 - Kernel functions which can be expressed as a dot product in some space satisfy the **Mercer's condition**
 - The Mercer's condition does not tell us how to construct $\Phi(\cdot)$ or even what the high dimensional space is.
 - see Burges' paper
- Advantages of kernel trick
 - no need to know $\Phi(\cdot)$
 - computations remain feasible even if the feature space has high dimensionality.

FLIP: Pattern Recognition (G.1)
46

46

Polynomial Kernel

- $K(\mathbf{x}, \mathbf{y}) = (\mathbf{x} \cdot \mathbf{y})^d$
 - * It can be shown for the case of polynomial kernels that the data is mapped to a space of dimension $h = \binom{p+d-1}{d}$ where p is the original dimensionality.
 - * Suppose $p=256$ and $d=4$, then $h=183,181,376$!!
 - * A dot product in the high dimensional space would require $O(h)$ computations while the kernel requires only $O(p)$ computations.

FLIP: Pattern Recognition (G.1)
47

47

Choice of Φ is not unique

Example: consider $\mathbf{x} \in \mathbb{R}^2$, $\Phi(\mathbf{x}) = \begin{pmatrix} x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{pmatrix} \in \mathbb{R}^3$, and $K(\mathbf{x}, \mathbf{y}) = (\mathbf{x} \cdot \mathbf{y})^2$

$$(\mathbf{x} \cdot \mathbf{y})^2 = (x_1y_1 + x_2y_2)^2$$

$$\Phi(\mathbf{x}) \cdot \Phi(\mathbf{y}) = x_1^2y_1^2 + 2x_1x_2y_1y_2 + x_2^2y_2^2 = (x_1y_1 + x_2y_2)^2$$

- Note that neither the mapping $\Phi(\cdot)$ nor the high dimensional space are unique.

$$\Phi(\mathbf{x}) = \frac{1}{\sqrt{2}} \begin{pmatrix} x_1^2 - x_2^2 \\ 2x_1x_2 \\ x_1^2 + x_2^2 \end{pmatrix} \in \mathbb{R}^3 \quad \text{or} \quad \Phi(\mathbf{x}) = \begin{pmatrix} x_1^2 \\ x_1x_2 \\ x_1x_2 \\ x_2^2 \end{pmatrix} \in \mathbb{R}^4$$

FLIP: Pattern Recognition (G.1)
48

48

Kernel functions

- By using different kernel functions, SVM implement a variety of learning machines, some of which coincide with classical architectures

polynomial: $K(x, x_k) = (x \cdot x_k)^d$

sigmoidal: $K(x, x_k) = \tanh(v_k(x \cdot x_k) + c_k)$
(corresponds to a two-layer sigmoidal neural network)

Gaussian: $K(x, x_k) = \exp\left(-\frac{\|x - x_k\|^2}{2\sigma_k^2}\right)$
(corresponds to a radial basis function (RBF) neural network)


49

Example

- Consider the XOR problem which is non-linear separable

(1,1) and (-1, -1) belong to ω_1

(1,-1) and (-1, 1) belong to ω_2



50

Example

- Consider the following mapping (among others)

$$y = \Phi(x) = \begin{pmatrix} x_1^2 \\ \sqrt{2}x_1 \\ \sqrt{2}x_2 \\ x_2^2 \\ 1 \end{pmatrix} \quad h=6$$

51

Example

- The above transformation maps x_k to a 6-dimensional space

$$y_1 = \Phi(x_1) = \begin{pmatrix} 1 \\ \sqrt{2} \\ \sqrt{2} \\ 1 \\ 1 \\ 1 \end{pmatrix} \quad y_3 = \Phi(x_3) = \begin{pmatrix} 1 \\ -\sqrt{2} \\ \sqrt{2} \\ -\sqrt{2} \\ 1 \\ 1 \end{pmatrix}$$

$$y_2 = \Phi(x_2) = \begin{pmatrix} 1 \\ \sqrt{2} \\ -\sqrt{2} \\ -\sqrt{2} \\ 1 \\ 1 \end{pmatrix} \quad y_4 = \Phi(x_4) = \begin{pmatrix} 1 \\ -\sqrt{2} \\ -\sqrt{2} \\ \sqrt{2} \\ 1 \\ 1 \end{pmatrix}$$

52

Example

- We seek to maximize

$$\sum_{k=1}^4 \lambda_k - \frac{1}{2} \sum_{k,j} \lambda_k \lambda_j z_k z_j \Phi(x'_j) \Phi(x_k)$$

subject to $\sum_{k=1}^4 z_k \lambda_k = 0, \lambda_k \geq 0, k = 1, 2, \dots, 4$

- The solution turns out to be:

$$\lambda_1 = \lambda_2 = \lambda_3 = \lambda_4 = \frac{1}{8}$$

- Since all $\lambda_k \neq 0$, all x_k are support vectors !

53

Example

- We now compute w

$$w = \sum_{k=1}^4 z_k \lambda_k \Phi(x_k) = \frac{1}{8} \begin{pmatrix} 1 \\ \sqrt{2} \\ \sqrt{2} \\ 1 \\ 1 \\ 1 \end{pmatrix} - \frac{1}{8} \begin{pmatrix} 1 \\ \sqrt{2} \\ -\sqrt{2} \\ 1 \\ 1 \\ 1 \end{pmatrix} + \frac{1}{8} \begin{pmatrix} 1 \\ -\sqrt{2} \\ \sqrt{2} \\ -\sqrt{2} \\ 1 \\ 1 \end{pmatrix} - \frac{1}{8} \begin{pmatrix} 1 \\ -\sqrt{2} \\ -\sqrt{2} \\ \sqrt{2} \\ 1 \\ 1 \end{pmatrix} = \frac{1}{2} \begin{pmatrix} 0 \\ \sqrt{2} \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$

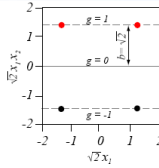
- The solution for w_0 can be determined using any support vector, e.g., x_1 :

$$w^T \Phi(x_1) + w_0 = z_1 \quad \text{or} \quad w_0 = 0$$

54

Example

- The margin b is computed as follows

$$w = \frac{1}{2} \begin{pmatrix} 0 \\ 0 \\ \sqrt{2} \\ 0 \\ 0 \\ 0 \end{pmatrix} \quad b = \frac{1}{\|w\|} = \sqrt{2}$$


TULIP Centre for Intelligent Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 55

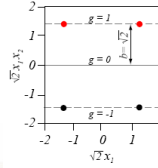
55

Example

- The discriminant function is as follows

$$g(x) = w^T \Phi(x) + w_0 = x_1 x_2$$

where we decide ω_1 if $g(x) > 0$ and ω_2 if $g(x) < 0$

$$w = \frac{1}{2} \begin{pmatrix} 0 \\ 0 \\ \sqrt{2} \\ 0 \\ 0 \\ 0 \end{pmatrix} \quad \Phi(x) = \begin{pmatrix} x_1^2 \\ \sqrt{2} x_1 x_2 \\ \sqrt{2} x_2 \\ x_2^2 \\ 1 \end{pmatrix}$$


TULIP Centre for Intelligent Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 56

56

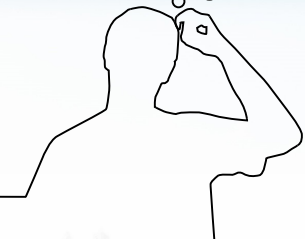
Summary of SVM

- SVM is based on exact optimization, not on approximate methods
 - i.e., global optimization method, no local optima
- Avoid overfitting in high dimensional spaces and generalize well using a small training set.
- Performance depends on the choice of the kernel and its parameters.
- Its complexity depends on the number of support vectors, not on the dimensionality of the transformed space.

TULIP Centre for Intelligent Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 57

57

Questions?



TULIP Centre for Intelligent Learning and Intelligent Processing FLIP: Pattern Recognition (G. Li) 59

59