

FUNDAMENTALS OF LEARNING AND INFORMATION PROCESSING

SESSION 17: STATISTICAL MACHINE LEARNING (VII)



Gang Li

Deakin University, Australia

2021-10-16

Regularized Loss Minimization	3
Regularized Loss Minimization (RLM).....	4
Stability	5
The Fitting-Stability Trade-off	8
Support Vector Machine	11
Binary Classification	12
Loss Function.....	13
Margin	14
Support Vector Machine (Hard-SVM).....	16
Support Vector Machine (Soft-SVM).....	17
Support Vectors	19
Representer Theorem	20
Kernel Trick	21
Embeddings into feature spaces	22
Dual Representation of Hypothesis	24
Kernel Trick	25
Mercer’s Condition.....	26
From Machine Learning to Deep Learning	27
Quiz	28

Table of Content

Regularized Loss Minimization

- Regularized Loss Minimization (RLM)
- Stability
- The Fitting-Stability Trade-off

Support Vector Machine

- Binary Classification
- Loss Function
- Margin
- Support Vector Machine (Hard-SVM)
- Support Vector Machine (Soft-SVM)
- Support Vectors
- Representer Theorem

Kernel Trick

- Embeddings into feature spaces
- Dual Representation of Hypothesis
- Kernel Trick
- Mercer’s Condition
- From Machine Learning to Deep Learning

Quiz

Regularized Loss Minimization

Regularized Loss Minimization (RLM)

Regularized Loss Minimization (RLM) is a learning paradigm in which we jointly minimize the empirical risk and a regularization function, which is a mapping $R : \mathcal{R}^d \mapsto \mathcal{R}$, the regularized loss minimization rule outputs a hypothesis which

$$\underset{\omega}{\operatorname{argmin}}(L_S(\omega) + R(\omega))$$

Tikhonov regularization is one popular regularization function: $R(\omega) = \lambda \|\omega\|^2$, where $\lambda > 0$ is a scalar, and the norm is the l_2 norm.

Notes.

- It is similar to SRM and MDL paradigm:
 - RLM and MDL** The “prior belief” of biasing to “short” vector in the \mathcal{H} .
 - RLM and SRM** We can define a sequence of hypothesis classes, $\mathcal{H}_1 \subset \mathcal{H}_2 \subset \mathcal{H}_3 \dots$, where $\mathcal{H}_i = \{\omega : \|\omega\| \leq i\}$. If the sample complexity of each \mathcal{H}_i depends on i , then the RLM is similar to SRM for this sequence of nested classes.
- Stabilizer:** Tikhonov regularization makes the learner stable w.r.t. small perturbation of the training set, which in turn leads to better generalization.

□

Stability

Given a training set $S = (z_1, \dots, z_{i-1}, z_i, z_{i+1}, \dots, z_m)$ and an additional example z' , let $S^{(i)}$ be the training set obtained by replacing $z_i \in S$ by z' , namely $S^{(i)} = (z_1, \dots, z_{i-1}, z', z_{i+1}, \dots, z_m)$ and let $U(m)$ be the uniform distribution over $[m]$. Let $\epsilon : \mathcal{N} \mapsto \mathcal{R}$ be a monotonically decreasing function. We say that a learning algorithm A is **On-Average-Replace-One-Stable** with rate $\epsilon(m)$ if every distribution \mathcal{D} :

$$\mathbb{E}_{(S, z') \sim \mathcal{D}^{m+1}, i \sim U(m)} [l(A(S^{(i)}), z_i) - l(A(S), z_i)] \leq \epsilon(m)$$

Notes.

- Informally: an algorithm A is **stable** if a small change of its input S will lead to a small change of its output hypothesis.
- Need to specify what is “**small change of input**” and what is “**small change of output**”.

□

Stable Rules Do Not Overfit

if A is *on-average-replace-one-stable* with rate $\epsilon(m)$ then

$$\mathbb{E}_{S \sim \mathcal{D}^m} [L_{\mathcal{D}}(A(S)) - L_S(A(S))] \leq \epsilon(m)$$

Proof.

- Since S and z' are both drawn i.i.d. from \mathcal{D} , we have that for every i

$$\mathbb{E}_S [L_{\mathcal{D}}(A(S))] = \mathbb{E}_{(S, z')} [l(A(S), z')] = \mathbb{E}_{(S, z')} [l(A(S^{(i)}), z_i)]$$

- On the other hand, we can write

$$\mathbb{E}_S [L_S(A(S))] = \mathbb{E}_{(S), i} [l(A(S), z_i)]$$

- The proof follows from the definition of *stability*.

□

Tikhonov Regularization as Stabilizer

Assume that the loss function is convex and ρ -Lipschitz. Then, the RLM rule with the regularizer $\lambda\|\omega\|^2$ is on-average-replace-one-stable with rate $\frac{2\rho^2}{\lambda m}$. It follows that:

👉
$$\mathbb{E}_{S \sim \mathcal{D}^m} [L_{\mathcal{D}}(A(S)) - L_S(A(S))] \leq \frac{2\rho^2}{\lambda m}$$

Similarly, for convex, β -smooth, and non-negative, the loss rate is $\frac{48\beta C}{\lambda m}$, with C is the upper bound on $\max_z l(\vec{0}, z)$, where $\omega = \vec{0}$.

Notes.

- The proof relies on the notion of strong convexity and is omitted here.

□

(None)-174ca9a (2021-10-16) – 7 / 31

The Fitting-Stability Trade-off

The expected risk of a learning algorithm A can be rewritten as

👉
$$\mathbb{E}_S [L_{\mathcal{D}}(A(S))] = \mathbb{E}_S [L_S(A(S))] + \mathbb{E}_S [L_{\mathcal{D}}(A(S)) - L_S(A(S))]$$

Notes.

- The first term is how good A fits the training set.
- The second term is the overfitting, and is bounded by the stability of A .
- λ controls the trade-off between above two terms.

□

Notes.

- Let A be the RLM rule.
- We saw (for convex-Lipschitz losses) $\mathbb{E}_{S \sim \mathcal{D}^m} [L_{\mathcal{D}}(A(S)) - L_S(A(S))] \leq \frac{2\rho^2}{\lambda m}$
- Fix some arbitrary vector ω^* , then $L_S(A(S)) \leq L_S(A(S)) + \lambda \|A(S)\|^2 \leq L_S(\omega^*) + \lambda \|\omega^*\|^2$.
- Taking expectation of both sides with respect to S and noting that $\mathbb{E}_S [L_S(\omega^*)] = L_{\mathcal{D}}(\omega^*)$, we obtain that $\mathbb{E}[L_S(A(S))] \leq L_{\mathcal{D}}(\omega^*) + \lambda \|\omega^*\|^2$.
- Therefore,

$$\mathbb{E}[L_{\mathcal{D}}(A(S))] \leq L_{\mathcal{D}}(\omega^*) + \lambda \|\omega^*\|^2 + \frac{2\rho^2}{\lambda m}$$

- The stability term decreases as λ increases, and the empirical risk increases with λ . So a trade-off is needed.

□

(None)-174ca9a (2021-10-16) – 8 / 31

The Regularization Path


👉 The RLM rule as a function of λ is $\omega(\lambda) = \operatorname{argmin}_{\omega} L_S(\omega) + \lambda \|\omega\|^2$. It can be seen as a *Pareto* objective: minimize both $L_S(\omega)$ and $\|\omega\|^2$.

How to choose λ .

Bound minimization choose λ according to the bound on $L_{\mathcal{D}}(\omega)$ usually far from optimal as the bound is the worst case.

Validation calculate several *Pareto* optimal points on the regularization path (by varying λ) and use validation set to choose the best one.

□



(None)-174ca9a (2021-10-16) – 9 / 31

Dimension vs. Norm Bounds

👉 The expected risk of a learning algorithm A can be rewritten as

$$E[L_{\mathcal{D}}(A(S))] \leq L_{\mathcal{D}}(\omega^*) + \lambda \|\omega^*\|^2 + \frac{2\rho^2}{\lambda m}$$

Notes.

- Previously in the course, when we learned d parameters, the *sample complexity* grew with d .
- Here, we learn d parameters but the *sample complexity* depends on the norm of $\|\omega\|$ and on the Lipschitzness/smoothness, rather than on d .
- Which approach is better depends on the properties of the distribution.

□

(None)-174ca9a (2021-10-16) – 10 / 31

Binary Classification

Consider a binary classification problem:

Hypothesis \mathcal{H}

the function set as:

$$h(x) = \begin{cases} 1 & \text{when } f(x) > 0 \\ -1 & \text{when } f(x) < 0 \end{cases}$$

👍

Loss Function

The number of times h get incorrect results on the sample.

$$L(h(x), y) = \sum_{i=1}^n l^{0-1}(h(x_i) \neq y_i) \cong \sum_{i=1}^n l(f(x_i), y_i)$$

Training by Optimization

Gradient descent is possible if both $h(x)$ and $f(x)$ are differentiable, otherwise difficult.

(None)-174ca9a (2021-10-16) – 12 / 31

Loss Function

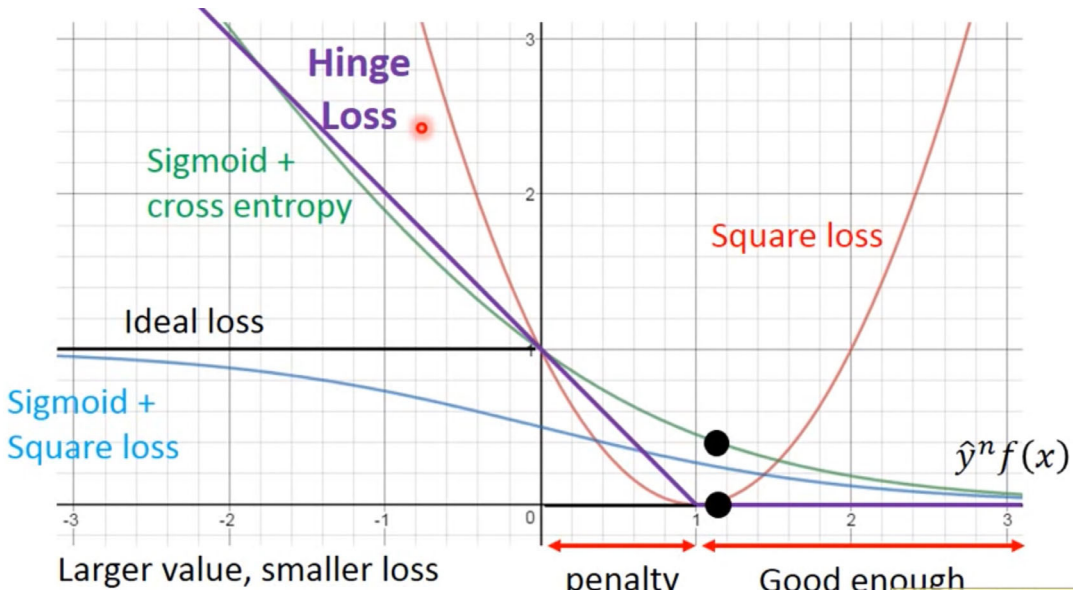
How to choose a differentiable function to approximate l^{0-1} ?

Loss Function

The number of times h get incorrect results on the sample.

👍

$$L(h(x), y) = \sum_{i=1}^n l^{0-1}(h(x_i) \neq y_i) \cong \sum_{i=1}^n l(f(x_i), y_i)$$



Squared Loss.

$$l(f(x_i), y_i) = (y_i f(x_i) - 1)^2$$

■

Intuitively, it wants to achieve:

$$f(x_i) = \begin{cases} 1 & \text{when } y_i = 1 \\ -1 & \text{when } y_i = -1 \end{cases}$$

■

It penalizes the very correct examples where $y_i f(x_i) \gg 1$

□

Sigmoid + Squared Loss.

$$l(f(x_i), y_i) = \sigma(y_i f(x_i)) - 1)^2 = \begin{cases} \sigma(f(x_i)) - 1)^2 & \text{when } y_i = 1 \\ \sigma(f(x_i))^2 & \text{when } y_i = -1 \end{cases}$$

■

It serves the purpose by achieving

$$\sigma(y_i f(x_i)) = \begin{cases} 1 & \text{when } y_i = 1 \\ 0 & \text{when } y_i = -1 \end{cases}$$

□

Sigmoid + Cross Entropy Loss.

$$l(f(x_i), y_i) = \ln(1 + e^{-y_i f(x_i)})$$

■

It achieve the cross entropy between two Bernoulli distributions: $(y_i, 1 - y_i)$ and $\sigma(f(x_i)), 1 - \sigma(f(x_i))$. Here we divide it by $\ln 2$ so that it is a surrogate loss function for l^{0-1} .

■

It serves the purpose.

□

Hinge Loss.

$$l(f(x_i), y_i) = \max(0, 1 - y_i f(x_i))$$

■

It achieves

$$f(x_i) = \begin{cases} \geq 1 & \text{when } y_i = 1 \\ \leq -1 & \text{when } y_i = -1 \end{cases}$$

■

It serves the purpose but different from *Sigmoid + Cross Entropy* loss.

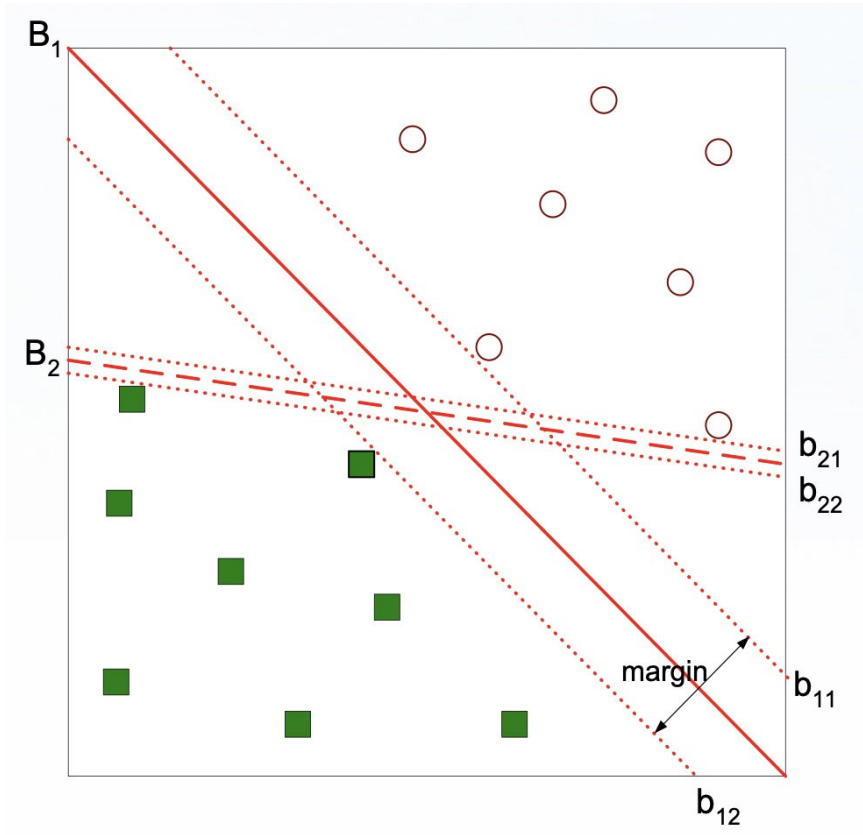
□

(None)-174ca9a (2021-10-16) – 13 / 31

6

Margin

Which separating hyperplane is better?




- Intuitively, solid red line is better.

Margin

Given hyperplane defined by $L = \{v : \langle \omega, v \rangle + b = 0\}$ and give a point x , the distance of x to L is

$d(x, L) = \min \|x - v\| : v \in L$

- If $\|\omega\| = 1$, then $d(x, L) = \|\langle \omega, x \rangle + b\|$



- Proof can be done easily.
- Some observation on the inner product:
 $\langle \omega, x \rangle = \|\omega\| \cdot \|x\| \cdot \cos(\theta)$

Support Vector Machine (Hard-SVM)

When the sample is linearly separable, we seek for the separating hyperplane with largest margin $\operatorname{argmax}_{(\omega,b):\|\omega\|=1} \min_{i\in m} \|\langle \omega, x_i \rangle + b\|$, subject to $\forall i, y_i(\langle \omega, x_i \rangle + b) > 0$.

■ equivalent to $\operatorname{argmax}_{(\omega,b):\|\omega\|=1} \min_{i\in m} y_i(\langle \omega, x_i \rangle + b)$

■ equivalent to $(\omega_0, b_0) = \operatorname{argmin}_{(\omega,b)} \|\omega\|^2$ subject to $\forall i, y_i(\langle \omega, x_i \rangle + b) \geq 1$.

■ the margin of $(\frac{\omega_0}{\|\omega_0\|}, \frac{b_0}{\|\omega_0\|})$ is $\frac{1}{\|\omega_0\|}$, and it is the maximal margin.

Notes.

- Margin is Scale Sensitive** The margin depends on the scale of the examples
- if (ω, b) separates $(x_1, y_1), \dots, (x_m, y_m)$ with margin γ , then it separates $(2x_1, y_1), \dots, (2x_m, y_m)$ with a margin of 2γ
- Margin of distribution** We say that \mathcal{D} is separable with a (γ, ρ) -margin if exists (ω^*, b^*) s.t. $\|\omega^*\| = 1$ and $\mathcal{D}(\{(x, y) : \|x\| \leq \rho \wedge y(\langle \omega^*, x \rangle + b^*) \geq \gamma\}) = 1$
- then its sample complexity is $m(\epsilon, \delta) \leq \frac{8}{\epsilon^2} 2(\rho/\gamma)^2 + \log(2/\delta)$
 - unlike the VC bounds, here the sample complexity depends on ρ/γ rather than d .

□

Support Vector Machine (Soft-SVM)

What if the sample is not linearly separable, we seek for the separating hyperplane with slack variable ϵ_n , minimizing the loss function L with RLM:

$$\operatorname{argmin}_{(\omega,b):\|\omega\|=1} L(\omega, S) = \sum_{i=1}^m l(f(x_i), y_i) + \lambda \|\omega\|^2 = \sum_{i=1}^m \epsilon_i + \lambda \|\omega\|^2$$

where $\epsilon_i = l^{hinge}(f(x_i), y_i) = \max(0, 1 - y_i f(x_i))$.

- the constraints are equivalent to $\epsilon_i \geq \begin{cases} 0 \\ 1 - y_i f(x_i) \end{cases}$
- this second one is equivalent to $y_i f(x_i) \geq 1 - \epsilon_i$

Notes.

- This is the popular SVM formulation, which can be solved by *quadratic programming*.
- As an optimization problem, it can also be solved by *gradient descendant*.

□

Support Vector Machine (Soft-SVM): Gradient Descendant

What if the sample is not linearly separable, we seek for the separating hyperplane with slack variable ϵ_n , minimizing the loss function L with RLM:

$$\operatorname{argmin}_{(\omega,b): \|\omega\|=1} L(\omega,S) = \sum_{i=1}^m l(f(x_i), y_i) + \lambda \|\omega\|^2$$

Gradient Descendant: $f(x_i) = \omega^T x_i$.

- Take partial derivatives to each component ω_j :

$$\frac{\partial L(f(x_i), y_i)}{\partial \omega_j} = \sum \frac{\partial l(f(x_i), y_i)}{\partial \omega_j} = \sum \frac{\partial l(f(x_i), y_i)}{f(x_i)} \frac{\partial f(x_i)}{\partial \omega_j}$$

Here we ignore the regularization term for simplicity.

- $\frac{\partial l(f(x_i), y_i)}{f(x_i)} = \begin{cases} -y_i & \text{if } y_i f(x_i) < 1 \\ 0 & \text{otherwise} \end{cases}$ while $\frac{\partial f(x_i)}{\partial \omega_j} = (x_j)_i$
- So $\frac{\partial L(f(x_i), y_i)}{\partial \omega_j} = \sum I(y_i f(x_i) < 1)(-y_i)x_i = \sum C_i(\omega_j)(x_j)_i$
- For ω_j , the gradient descendant updating rule is $\omega_j = \omega_j - \eta \sum C_i(\omega_j)(x_j)_i$, in the vector form: $\omega = \omega - \eta \sum C_i(\omega)x_i$.


□

Support Vectors

A **separating hyperplane** is defined by (ω, b) subject to: $\forall i, y_i(\langle \omega, x_i \rangle + b) > 0$. The margin of a separating hyperplane is the distance of the closest example to it:

$$\min_i \|\langle \omega, x_i \rangle + b\|$$

Those closest examples are called **support vectors**.



- From the gradient descendant method, we can see that when $\omega = \vec{0}$, $\omega^* = \sum_i \alpha_i^* x_i$ is a linear combination of examples.
- α^* may be sparse, and those x_i with non-zero α_i^* are **support vectors**.
 - For *Hinge loss*, α^* is usually sparse.
 - For *logistic regression* or *cross entropy*, α^* is usually non-zero

Representer Theorem

Assume that ψ is a mapping from \mathcal{X} to a Hilbert space (a feature space), then the SVM optimization is an instance of the following problem:

$$\operatorname{argmin}_{\omega} (f(\langle \omega, \psi(x_1) \rangle, \dots, \langle \omega, \psi(x_m) \rangle) + R(\|\omega\|))$$

where $f : \mathcal{R}^m \mapsto \mathcal{R}$ is an arbitrary function. $R : \mathcal{R}_+ \mapsto \mathcal{R}$ is a monotonically no-decreasing function, such as $\lambda \|\omega\|^2$. Then $\exists \alpha \in \mathcal{R}^m$ such that $\omega^* = \sum_{i=1}^m \alpha_i \psi(x_i)$.

Intuition.

- Because ω^* is an element of a Hilbert space, so $\omega^* = \sum_{i=1}^m \alpha_i \psi(x_i) + u$, where $u \perp \psi(x_i) \ \forall x_i$.
- Set $\omega = \omega^* - u$, observe that $\omega^* = \omega + u$, we have $\|\omega^*\|^2 = \|\omega\|^2 + u^2$ and $\forall i \ \langle \omega, \psi(x_i) \rangle = \langle \omega^*, \psi(x_i) \rangle$.
- Hence the objective at ω equals the objective at ω^* minus $\lambda \alpha$. By optimality of ω^* , u must be zero.

□

Implications.

- By representer theorem, the optimal solution can be written as $\omega^* = \sum_{i=1}^m \alpha_i \psi(x_i)$
- Denote by \mathcal{G} the Gram matrix s.t. $\mathcal{G}_{i,j} = \langle \psi(x_i), \psi(x_j) \rangle$, we have:

$$\langle \omega, \psi(x_i) \rangle = \langle \sum_{i=1}^m \alpha_i \psi(x_i), \psi(x_i) \rangle = \sum_{i=1}^m \alpha_i \langle \psi(x_i), \psi(x_i) \rangle = (\mathcal{G} \alpha)_i, \quad \forall i$$

- Also $\|\omega\|^2 = \alpha^T \mathcal{G} \alpha$. Hence, the optimisation task can be written as:

$$\operatorname{argmin}_{\alpha \in \mathcal{R}^m} (f(\mathcal{G} \alpha) + \lambda \alpha^T \mathcal{G} \alpha)$$

□

(None)-174ca9a (2021-10-16) – 20 / 31

Kernel Trick

21 / 31

Embeddings into feature spaces


What if the sample S is not linear separable?

Notes.

- The following sample in \mathcal{R}^1 is not separable by half-spaces



- It is separable in \mathcal{R}^2 by half-spaces if we map $x \mapsto (x, x^2)$



□

(None)-174ca9a (2021-10-16) – 22 / 31

Embeddings into feature spaces

Define a mapping function $\psi : \mathcal{X} \mapsto \mathcal{F}$, where the feature space \mathcal{F} is a subset of Hilbert space. Then the training of half-space is done over

$$\{(\psi(x_1), y_1), \dots, (\psi(x_m), y_m)\}$$

Notes.

- How to choose ψ ?
 - ◆ In general, this requires prior knowledge.
 - ◆ There are some generic mappings that enrich the class of half-spaces, e.g. polynomial mappings.
- If F is high dimensional we face
 - statistical challenge** can be tackled using margin
 - computational challenge** can be tackled using kernels

□

Dual Representation of Hypothesis

We know that $\omega^* = \sum_i \alpha_i^* x_i = \vec{\alpha}^* \alpha^*$ is a linear combination of x_i . Accordingly, the hypothesis $h(x)$ can be written as

$$h(x) = \omega \cdot \vec{x} = \vec{\alpha}(\alpha \vec{x})^T = \alpha^T \vec{x}^T \vec{x} = \sum_i \alpha_i (x_i \cdot x_i) = \sum_i \alpha_i \mathcal{K}(x_i \cdot x_i)$$

Notes.

- In this representation, the training of a hypothesis is equivalent to find $\vec{\alpha}^* = \{\alpha_1, \dots, \alpha_m\}$, minimizing the loss function L .
- $$L = \sum_{i=1}^m l(\sum_{j=1}^m \alpha_j \mathcal{K}(\vec{x}_j, \vec{x}_i), y_i)$$
- The Kernel Trick $\mathcal{K}(\vec{x}_j, \vec{x}_i)$:
 - ◆ We don't really need to know vectors \vec{x}_j and \vec{x}_i .
 - ◆ We only need to know the inner product.
- In the mapped feature space, it is then $\mathcal{K}(x_j, x_i) = \langle \psi(\vec{x}_j), \psi(\vec{x}_i) \rangle$.

□

Kernel Trick

A kernel function for a mapping ψ is a function that implements inner product in the feature space, namely,

$$\mathcal{K}(x, y) = \langle \psi(x), \psi(y) \rangle$$

Polynomial Kernel. The k degree polynomial kernel is defined to be

$$\mathcal{K}(x, y) = (1 + \langle \psi(x), \psi(y) \rangle)^k$$

- Since ψ contains all the monomials up to degree k , a half space over the range of ψ corresponds to a polynomial predictor of degree k over the original space.
- Observe that calculating $\mathcal{K}(x, y)$ takes $O(n)$ time while the dimension of $\psi(x)$ is nk
- Consider mapping from two dimensional space to three dimensional space: $x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$ and $\psi(x) = \begin{bmatrix} x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{bmatrix}$
- $\mathcal{K}(x, y) = \langle \psi(x), \psi(y) \rangle = \begin{bmatrix} x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{bmatrix} \cdot \begin{bmatrix} y_1^2 \\ \sqrt{2}y_1y_2 \\ y_2^2 \end{bmatrix} = x_1^2y_1^2 + 2x_1x_2y_1y_2 + x_2^2y_2^2 = (x \cdot y)^2$

□

Gaussian kernel or “Radial Basis Function (RBF)” kernel. It is defined to be

$$\mathcal{K}(x, y) = e^{-\frac{\|x-y\|^2}{2\sigma}}$$

- Let the original instance space be R and consider the mapping ρ where for each non-negative integer $n \geq 0$ there exists an element $\psi_n(x)$ which equals to $\frac{1}{\sqrt{n!}}e^{-\frac{x^2}{2\sigma}}x^n$.
- ψ can have infinite dimension:

$$\mathcal{K}(x, y) = \langle \psi(x), \psi(y) \rangle = \sum_{i=1}^n \frac{1}{\sqrt{n!}}e^{-\frac{x^2}{2\sigma}}y^n \frac{1}{\sqrt{n!}}e^{-\frac{y^2}{2\sigma}}y^n = e^{-\frac{\|x-y\|^2}{2\sigma}}$$


- It can learn any polynomial function.

□

Sigmoid kernel. It is defined to be

$$\mathcal{K}(x, y) = \tanh(\langle x, y \rangle)$$

- When using the sigmoid kernel, it is actually working as $h(x) = \sum_i \alpha_i \tanh(x_i, x)$.
- It can be considered as one neural network with one hidden layer.



- ◆ The weight of each neuron is an example x_i .
- ◆ The number of support vectors is the number of neurons.

□

Mercer’s Condition

A symmetric function $\mathcal{K} : \mathcal{X} \times \mathcal{X} \mapsto \mathcal{R}$ implements an inner product in some Hilbert space if and only if it is positive semi-definite; namely $\forall v_i$, the Gram matrix, $\mathcal{G}(i, j) = \mathcal{K}(x_i, x_j)$, is a positive semidefinite matrix.

Notes.


- It can be learned or designed by prior knowledge.

□


From Machine Learning to Deep Learning

Notes.

- SVM is a feature mapping followed by linear classifier (half-space): SVM kernel is learnable, but not as perfectly done as in ANN.



- Deep learning is feature transformations + linear classifier (half-space).



□

(None)-174ca9a (2021-10-16) – 27 / 31

Quiz

28 / 31

SGD with Projection Step



A supermarket manager would like to learn which of his customers have babies on the basis of their shopping carts. Specifically, he sampled i.i.d. customers, where for customer i , let $x_i \subset \{1, \dots, d\}$ denote the subset of items the customer bought, and let $y_i \in \{1, -1\}$ be the label indicating whether this customer has a baby. As prior knowledge, the manager knows that there are k items such that the label is determined to be 1 iff the customer bought at least one of these k items. Of course, the identity of these k items is not known (otherwise, there was nothing to learn). In addition, according to the store regulation, each customer can buy at most s items.

- Help the manager to design a learning algorithm such that both its time complexity and its sample complexity are polynomial in s , k , and $\frac{1}{\epsilon}$.




(None)-174ca9a (2021-10-16) – 29 / 31

Questions?

Contact Information

Associate Professor **GANG LI**
School of Information Technology
Deakin University
Geelong, Victoria 3216, Australia



-  GANGLI@TULIP.ORG.AU
-  OPEN RESOURCES OF TULIP-LAB
-  TEAM FOR UNIVERSAL LEARNING AND INTELLIGENT PROCESSING